



# Modeling Relational Data with Cosmos DB

Steve Faulkner  
@southpolesteve

**AT THE  
SOUTH  
POLE!**





Microsoft



# Azure Cosmos





# Microsoft Azure Cosmos JavaScript SDK

This project provides JavaScript & Node.js SDK library for [SQL API](#) of [Azure Cosmos Database Service](#). This project also includes samples, tools, and utilities.

npm@latest v3.3.0 Azure Pipelines succeeded

```
// JavaScript
const { CosmosClient } = require("@azure/cosmos");

const endpoint = "https://your-account.documents.azure.com"; // Add your endpoint
const key = "[database account masterkey]"; // Add the masterkey of the endpoint
const client = new CosmosClient({ endpoint, key });

const databaseDefinition = { id: "sample database" };
const collectionDefinition = { id: "sample collection" };
const documentDefinition = { id: "hello world doc", content: "Hello World!" };

async function helloCosmos() {
  const { database } = await client.databases.create(databaseDefinition);
  console.log("created database");

  const { container } = await database.containers.create(collectionDefinition);
  console.log("created collection");

  const { resource } = await container.items.create(documentDefinition);
```

# Modeling Relational Data with Cosmos DB

# A brief history of databases



1974

SQL Invented

1981

Price per GB?

\$700,000

2017

Price per GB?

\$0.025



SQL - Storage Optimized  
NoSQL - Speed Optimized



# Azure Cosmos



# Azure Cosmos DB



Table

SQL

SQL



JavaScript

{LEAF}

API for MongoDB



Gremlin



Cassandra

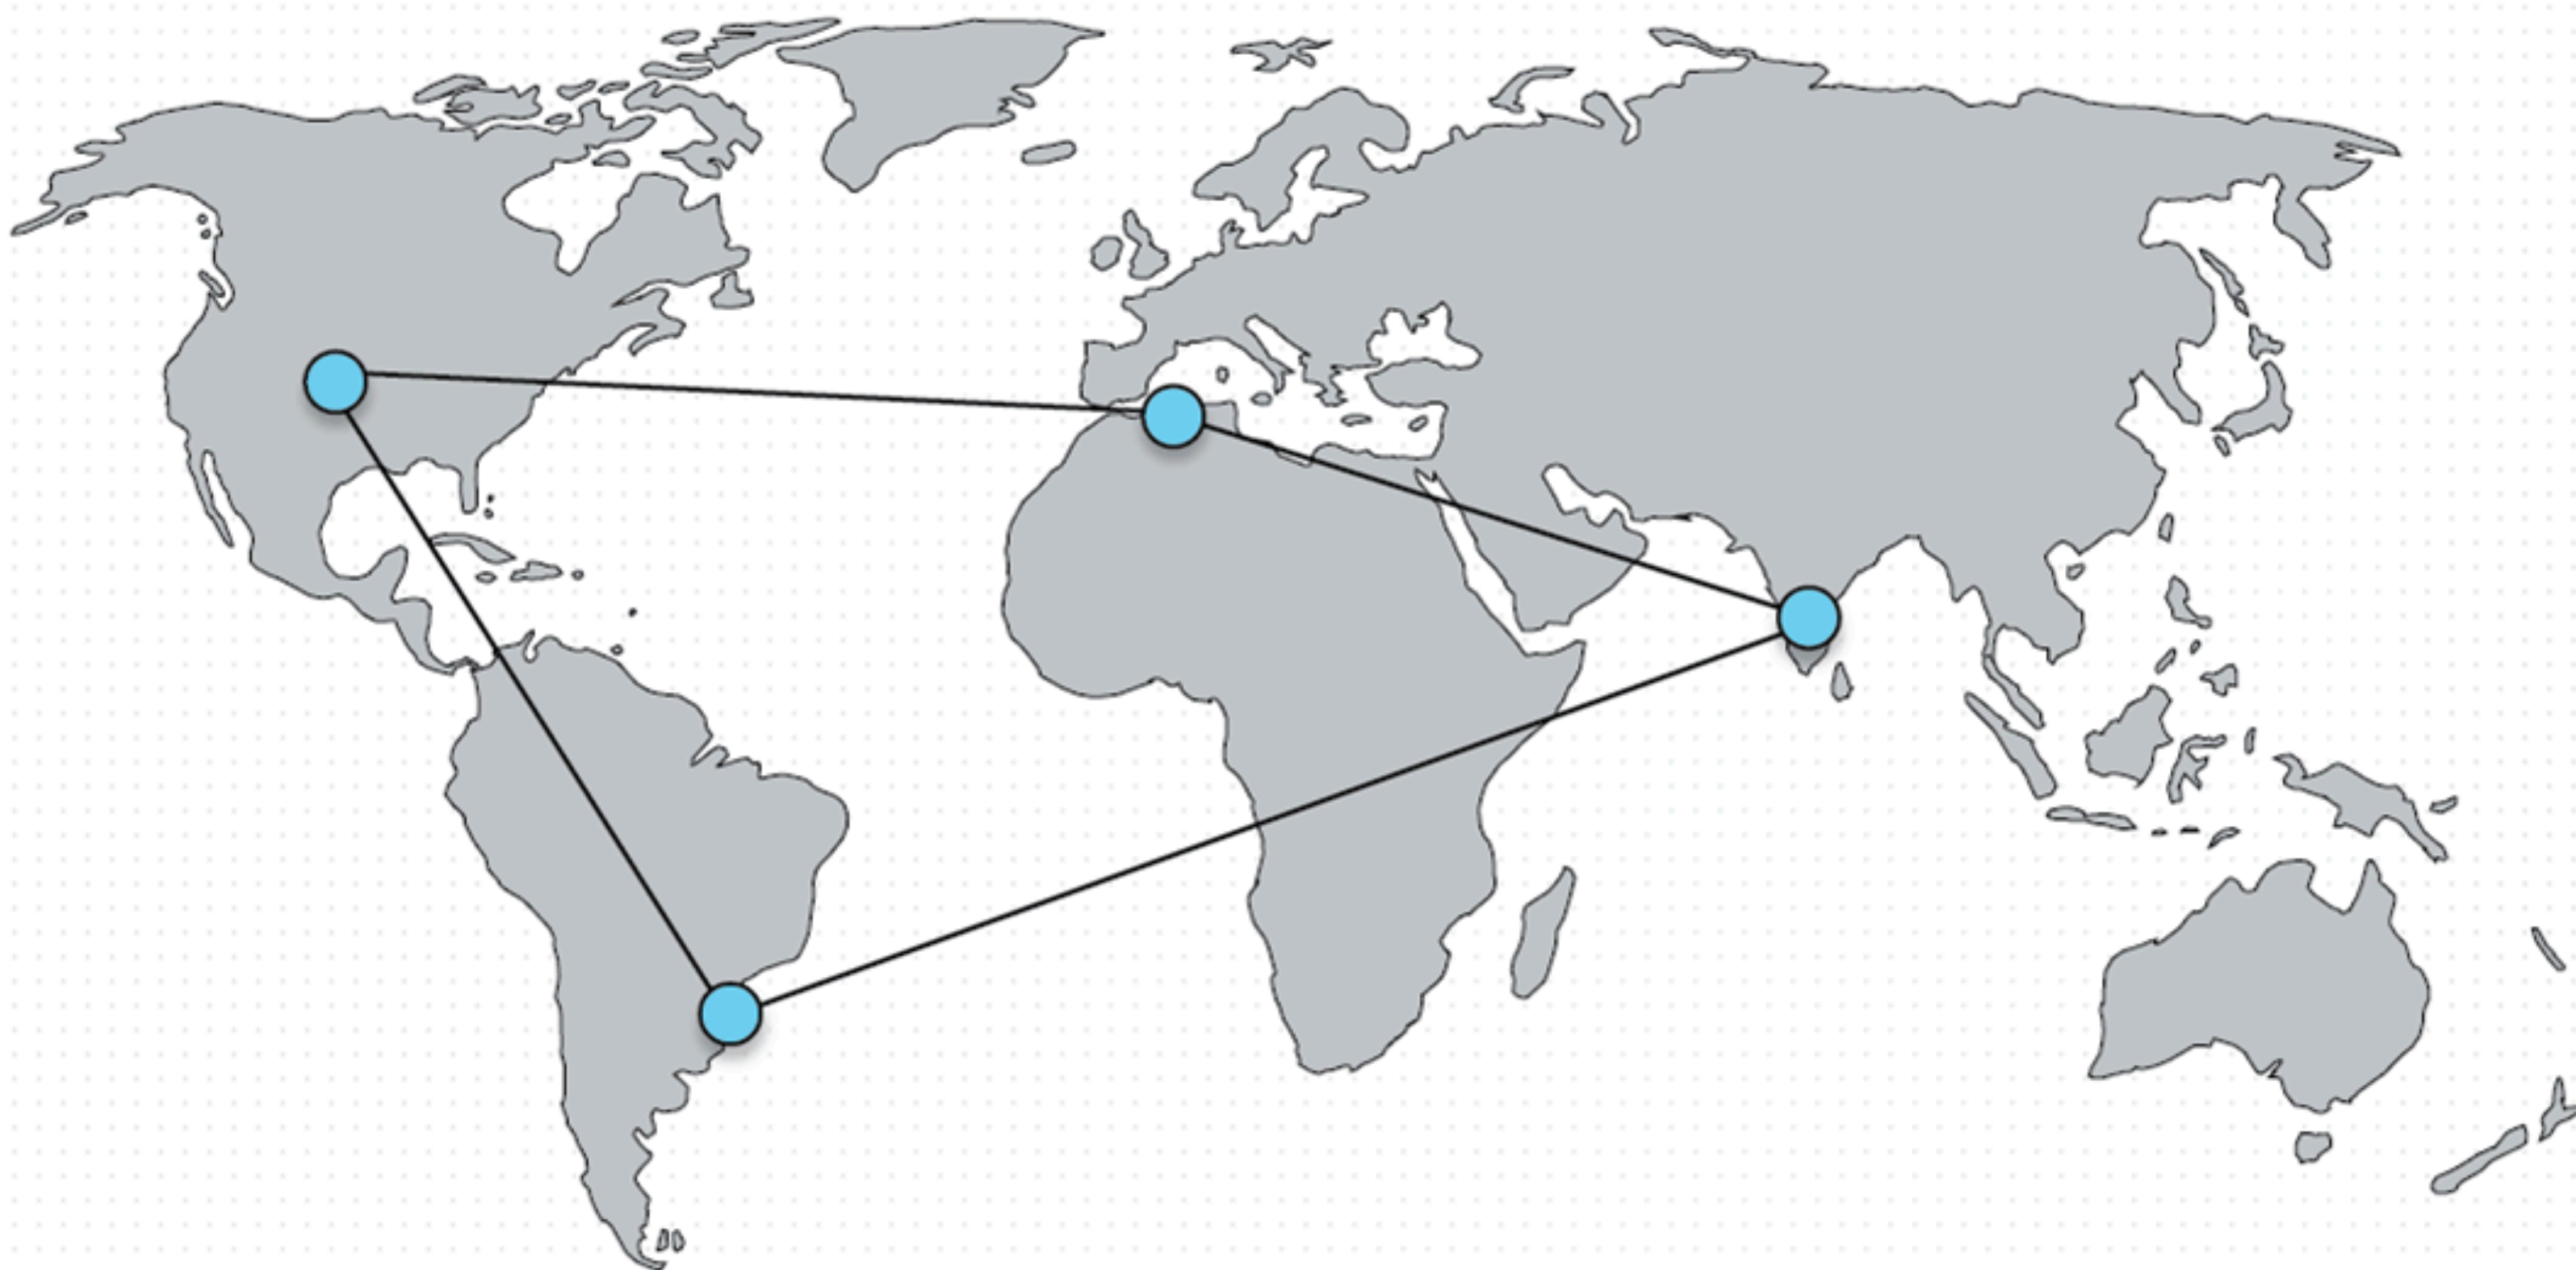


Spark




ETCD

...more APIs  
coming




Key-Value




Column-Family



Documents



Graph



Global distribution

Elastic scale-out

Guaranteed low latency

Five consistency models

Comprehensive SLAs

Read/Write JSON  
Massive scale+perf  
No schema  
Partitioned / Scale out  
SQL API  
Everything indexed by default

# Multiple APIs



SQL  
Tables  
Cassandra  
Mongo  
Gremlin

✨ Jupyter Notebooks ✨

✨ Spark ✨

✨ etcd ✨

Let's talk about  
“SQL”

SQL API

Not a SQL Database

# Queries



```
SELECT * from c
```

```
SELECT c.id from c
```

```
SELECT * from c  
WHERE c.dueAt > today
```

```
SELECT * from c  
WHERE c.priority > 10  
ORDER BY c.dueAt
```

SELECT, WHERE, DISTINCT,  
BETWEEN, IN, TOP, ORDER BY,  
COUNT, MAX, MIN, AVG



✨ Multi ORDER BY ✨

✨ OFFSET LIMIT ✨

✨ DISTINCT ✨




✨ Correlated Subqueries ✨

# NoSQL Gotchas


Scans

# Cross Partition Queries


# Partitioning



hash(partitionKey)



hash(partitionKey)







id + partitionKey

```
SELECT * from c WHERE  
c.partitionKey = "foo"
```

hash("foo")



hash("foo")




```
SELECT * from c WHERE  
c.notPartitionKey = "bar"
```

client



client







RECAP

SQL API != SQL


♥ partition key ♥


Avoid Cross Partition

# Relational Data

# Blog Example

Users  
Posts  
Comments





Users


PK: id



Posts



Comments




PK: id



PK: userId







PK: id




PK: userId



PK: postId





PK: id



PK: userId




PK: postId

# Cross Partition Queries

Posts  
pk: userId

```
Post.find(1)
```



Comments

pk: postId



Comment.find(1)

Comments.where(user: 1)


# SQL Thinking

Database schema

Normalized data

Rigid consistency

Many tables



PK: id



PK: userId



PK: postId

# NoSQL Thinking

Application schema  
Denormalized data  
Flexible consistency  
Single container



Data

PK:  $pk$

Post.find(1)



# Post.create()

```
{
```

```
  id: "post-1"
```

```
  pk: "post-1"
```

```
  title: "Cosmos is great!"
```

```
}
```

# Post.find(1)

{

id: "post-1"

pk: "post-1"

title: "Cosmos is great!"

}

```
User.find(1).posts()
```

# Post.create()

```
{  
  id: "post-2"  
  pk: "post-2"  
  authorId: "1"  
}
```

```
{  
  id: "post-2"  
  pk: "user-1-posts"  
}
```

```
{  
  id: "post-2"  
  pk: "user-1-posts"  
}
```

```
{  
  id: "post-3"  
  pk: "user-1-posts"  
}
```

```
{  
  id: "post-4"  
  pk: "user-1-posts"  
}
```


```
User.find(1).posts()
```

```
Select c.id from c WHERE  
c.pk = "user-1-posts"
```

```
User.find(1).posts()
```

```
Select c.id from c WHERE  
c.pk = "user-1-posts"
```

hash("user-1-posts")





hash("user-1-posts")



# User.find(1).posts()

```
[  
  { id: "post-2" },  
  { id: "post-3" },  
  { id: "post-4" },  
  ...  
]
```

```
User.find(1)  
.posts({ state: draft })
```

# Posts.create()

```
{  
  id: "post-2"  
  pk: "post-2"  
  authorId: "user-1"  
  state: "draft"  
}
```

```
{  
  id: "post-2"  
  pk: "user-1-posts"  
  state: "draft"  
}
```

# Posts.create()

```
{  
  id: "post-2"  
  pk: "post-2"  
  authorId: "user-1"  
  state: "draft"  
}
```

```
{  
  id: "post-2"  
  pk: "user-1-posts"  
  state: "draft"  
}
```

```
Select c.id from c  
WHERE c.pk = "user-1-posts"  
AND c.state = "draft"
```

# Posts.update()

```
{  
  id: "post-2"  
  pk: "post-2"  
  authorId: "user-1"  
  state: "draft"  
}
```

```
{  
  id: "post-2"  
  pk: "user-1-posts"  
  state: "draft"  
}
```

# Posts.update()

```
{  
  id: "post-2"  
  pk: "post-2"  
  authorId: "user-1"  
  state: "published"  
}
```

```
{  
  id: "post-2"  
  pk: "user-1-posts"  
  state: "published"  
}
```



User.findByEmail()

# User.create

```
{  
  id: "2"  
  pk: "2"  
  email: "steve@microsoft.com"  
}  
  
{  
  id: "2"  
  pk: "user-email-steve@microsoft.com"  
}
```

# User.findByEmail()

```
Select c.id from c WHERE  
c.email = "steve@microsoft.com"
```

# User.findByEmail()

Select c.id from c WHERE

c.pk = "user-emails-steve@microsoft.com"

# Advanced Techniques

# Update w/ Change Feed + Azure Functions

# Materialized Documents

# Split Containers



# Critical Takeaways

Single Container

Best partition key?  
“partitionKey”

Top Level Entities:

$id == pk$

Use partition key to  
index relationships and  
attributes

Optimize as data size grows  
and access patterns evolve

# Questions?

@southpolesteve

askcosmosdb@microsoft.com